Accelerating IBM Rational ClearCase Snapshot Views

Derek Lau
eBay
dlau@ebay.com

Let's build a smarter planet.

A Word of Warning

- Uncharted territory ahead!
- Not an IBM supported configuration
- You should be comfortable writing/debugging your own tools

The Problem

- Snapshot views load each load rule SEQUENTIALLY
  - Load <loadrule1>
  - Then load <loadrule2>
  - Then load <loadrule3>
  - etc

- There is only ONE single-threaded view_server process for each view
  - Unused CPU capacity, especially with multiple core CPUs
  - Unused network bandwidth

The Solution

- Create X “child” views
- Each child view will process a subset of the parent’s load rules
- Sample child view config spec

  element * CHECKEDOUT

  element * .../m_dlau/LATEST

  element * /main/LATEST -mkbranch m_dlau

  load /TeamSCM/ClearCase

- Combine the workspace, config spec, .compiled_spec, db into the parent view

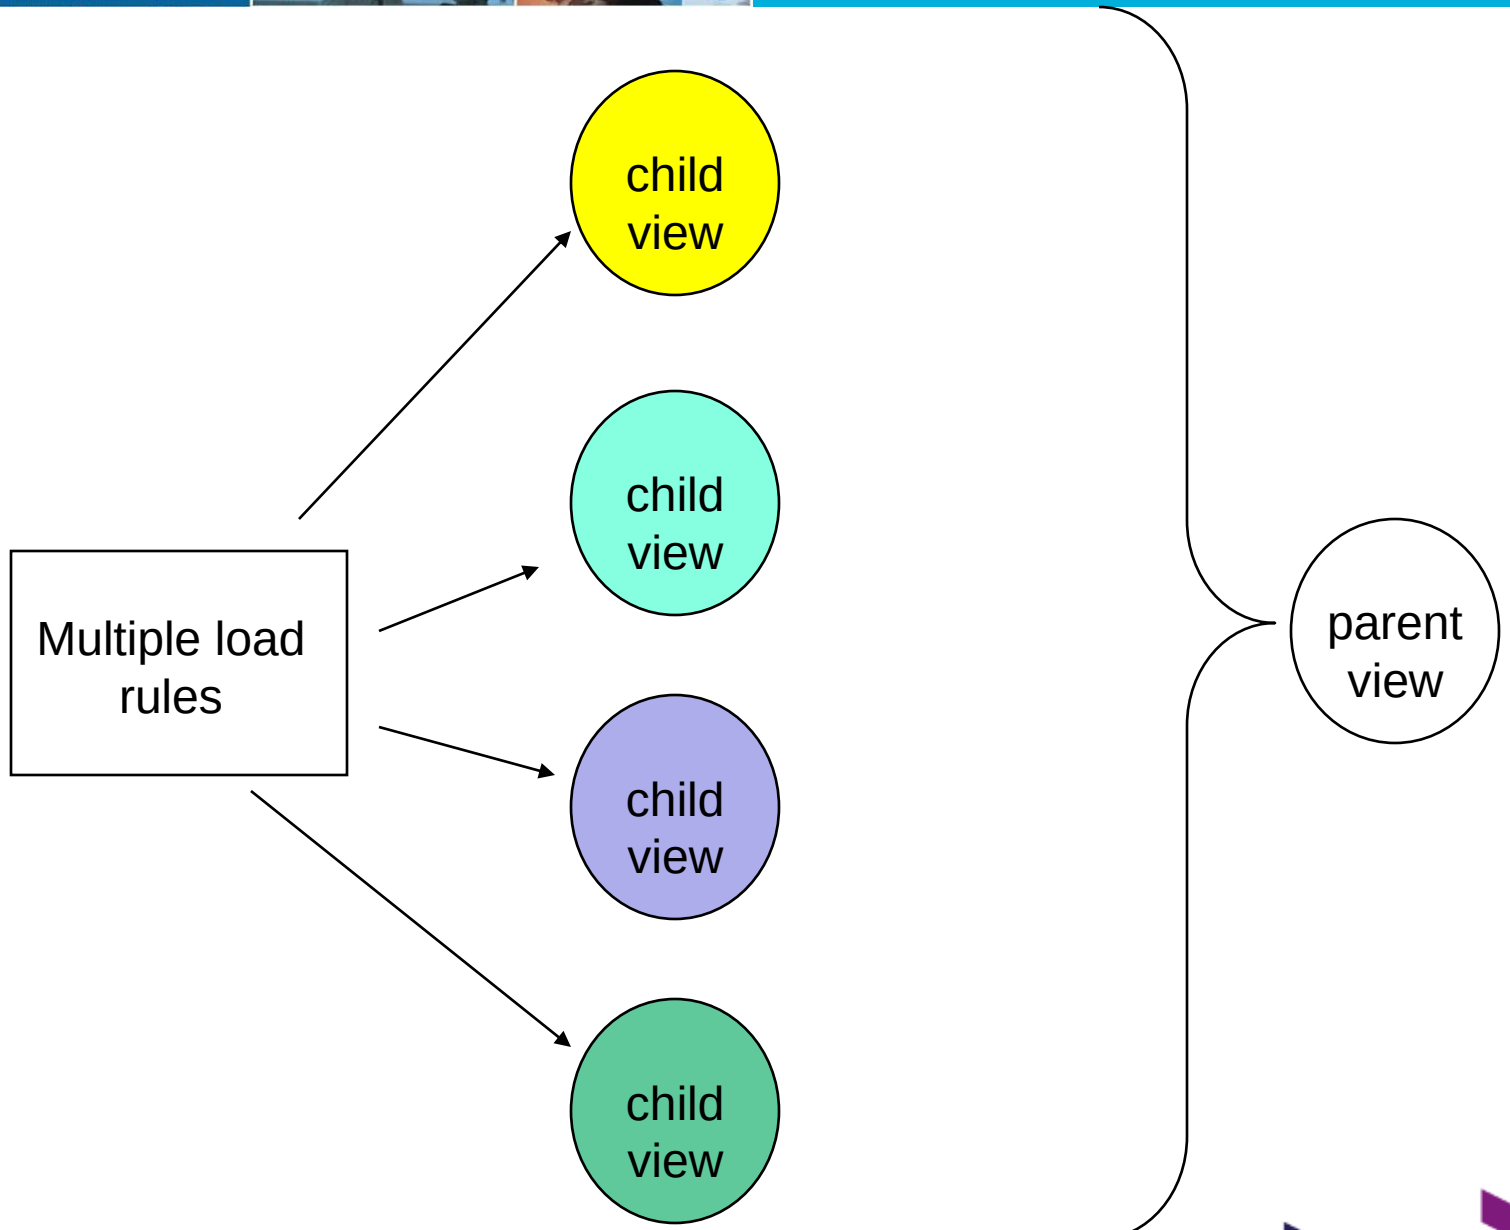

A Typical Snapshot View

- Typical snapshot view config spec

  element * CHECKEDOUT

  element * .../m_dlau/LATEST

  element * /main/LATEST -mkbranch m_dlau

  load /TeamSCM/ClearCase

  load /TeamSCM/Admin

  load /ICE/ICEConfig/scripts

  ... many other load rules ...

Technical Overview of a Snapshot View (1 / 3)

- Snapshot view workspace
  - identified by a view.dat file in the workspace directory root
- View storage (.vws or .stg)
  - config_spec
  - .compiled_spec
  - db directory
  - ... and other files
- view_server process
  - communicates with the VOB via vobrpc_server process
  - single threaded by design
  - updates the view database continuously during a load

Technical Overview of a Snapshot View (2 / 3)

- SENTENCE_1
- SENTENCE_2

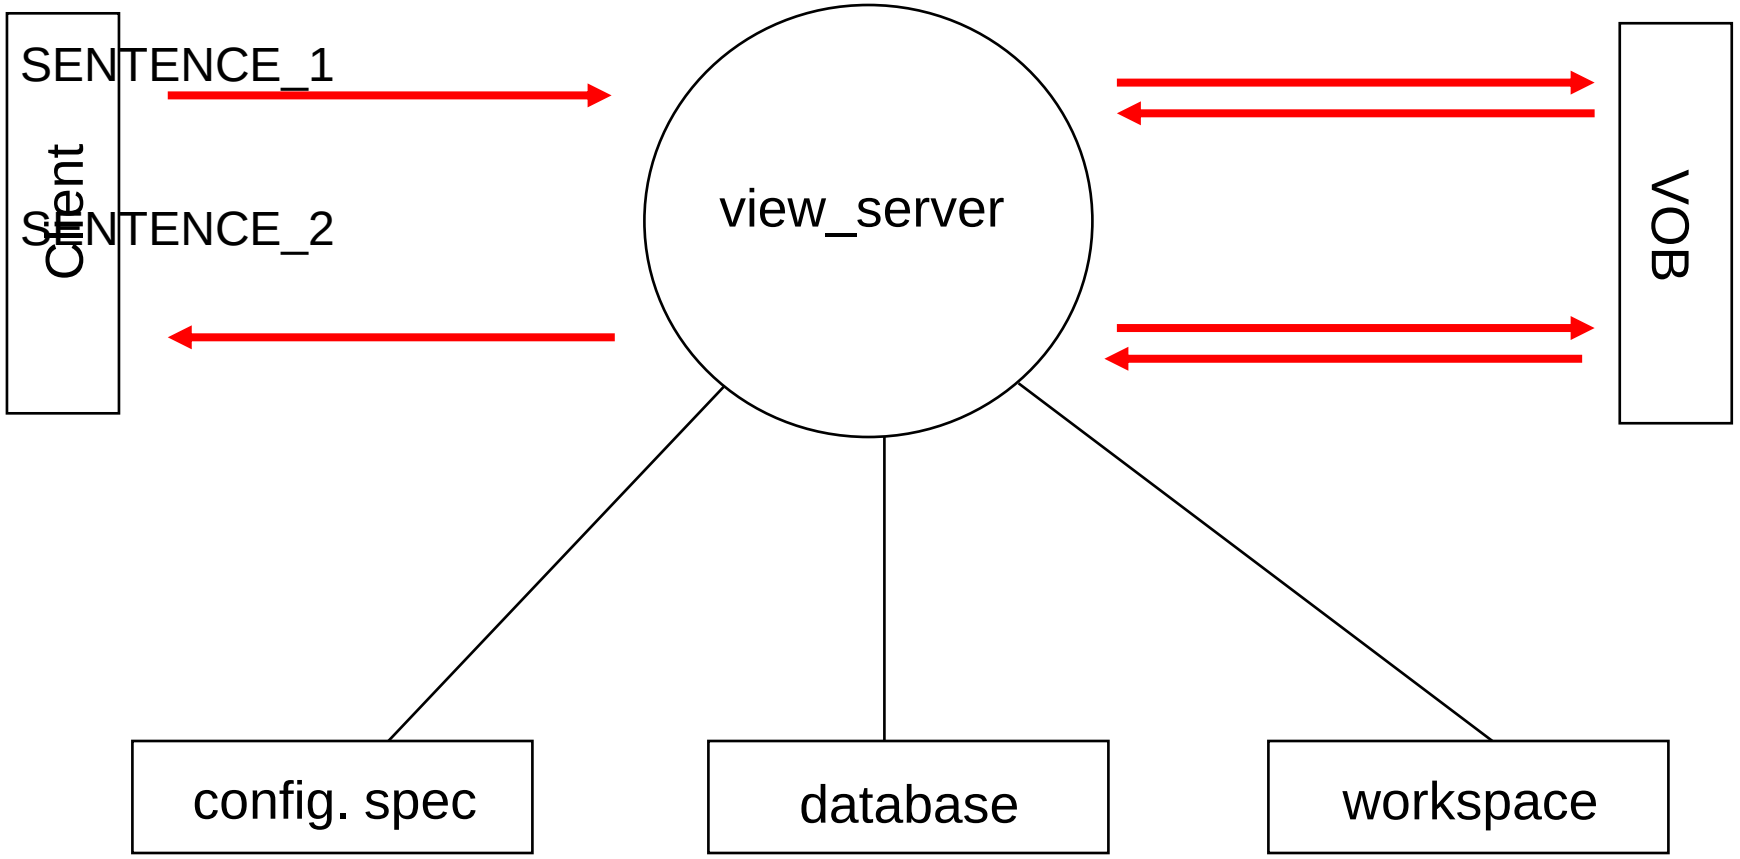

Technical Overview of a Snapshot View (3 / 3)

- What happens when you do setcs/edcs in a snapshot view?

- config_spec file compiled by view_server process to .compiled_spec
- each load rule is processed sequentially
  - view_server determines what version of this file/directory by parsing compiled_spec
  - view_server downloads one file
  - view_server update the view's database
  - view_server process next file

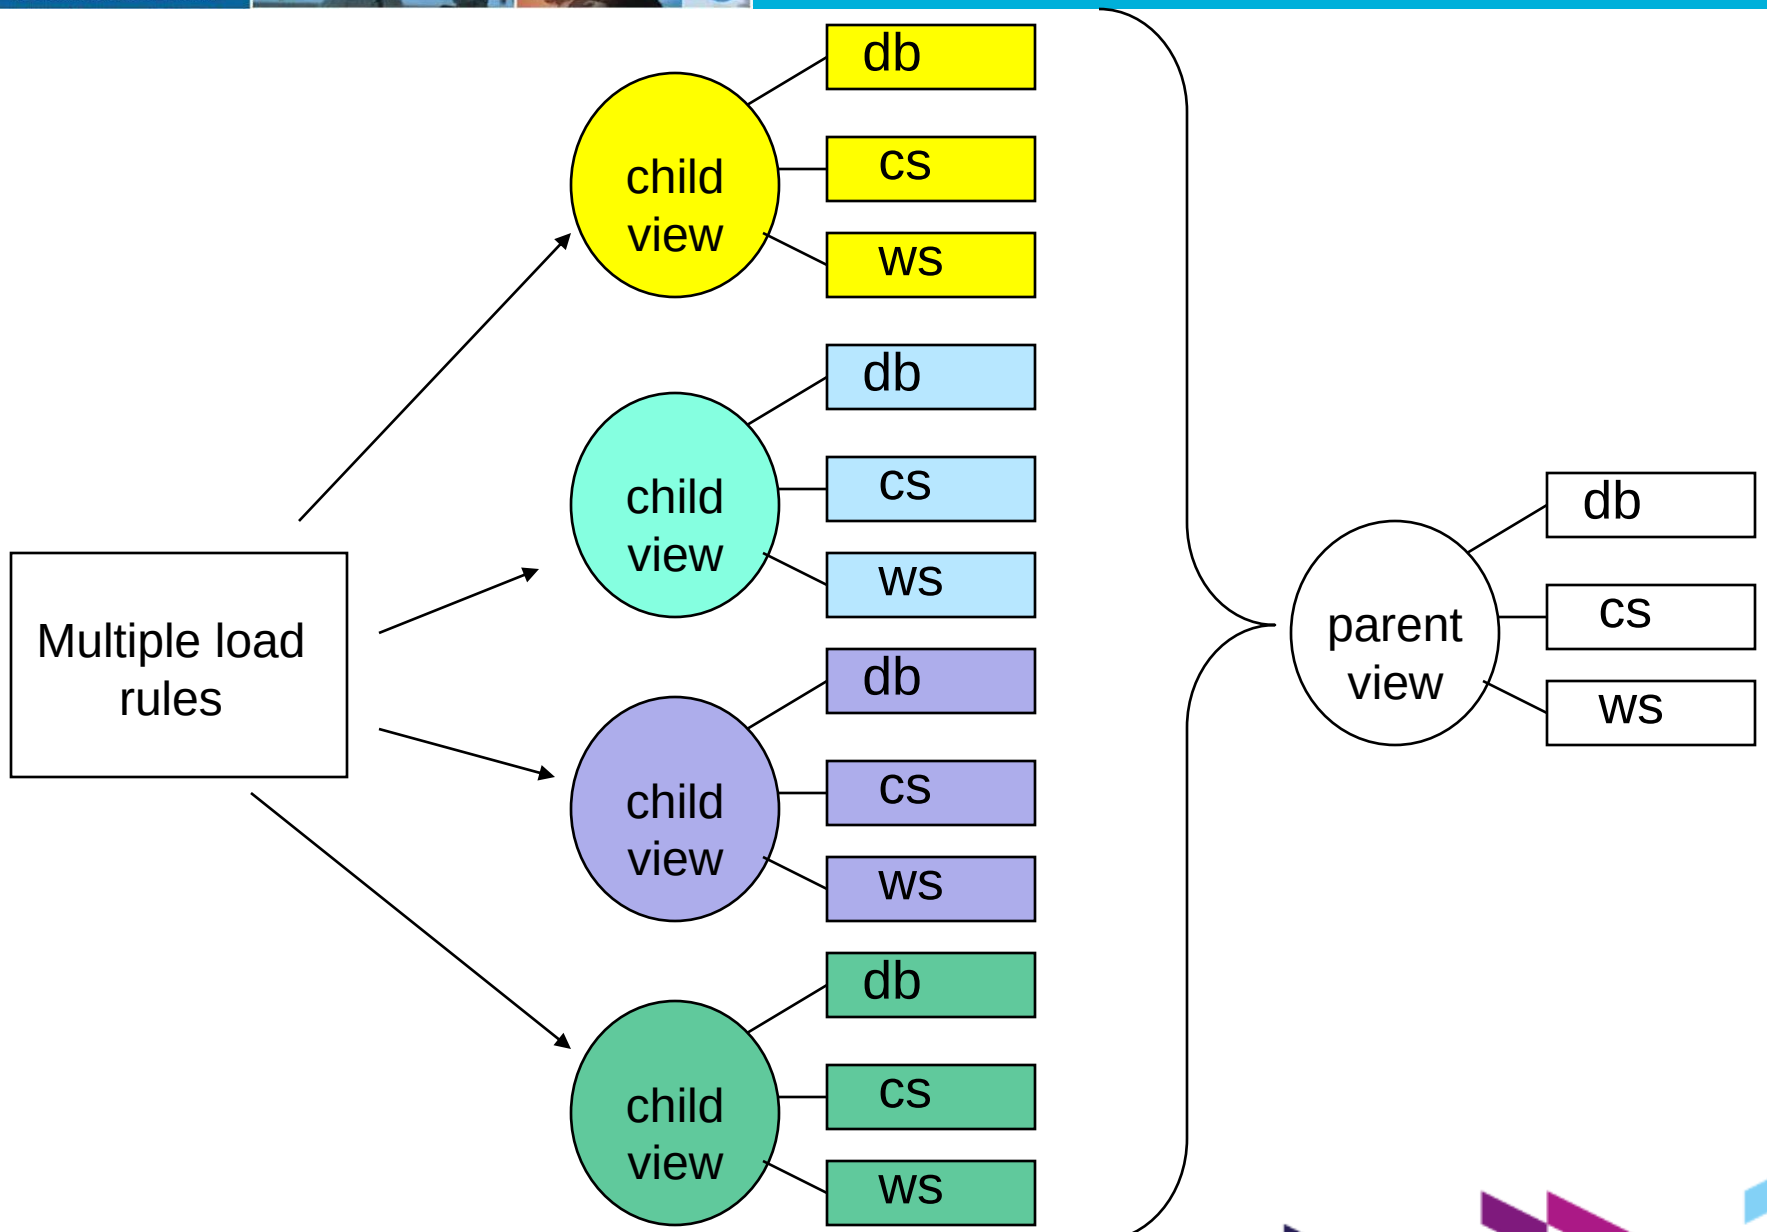

Reassembling the View Workspace

- Copy the files from the child workspaces into the parent workspace
- If we are only doing a build, we are done!
- If we need to do ClearCase operations (eg checkin, checkout), a few more steps...

Reassembling the config_spec file

- Simple human readable text file format

```
element * CHECKEDOUT
element * .../m_dlau/LATEST
element * /main/LATEST -mkbranch m_dlau
load /TeamSCM/ClearCase
load /TeamSCM/Admin
load /ICE/ICEConfig/scripts
```

Reassembling the .compiled_spec file (1 / 5)

- Compiled spec file for our example config spec looks like this:

```
12
dfe36c75.23684e37.ab97.9f:1f:b4:6d:7c:97
6
1 1 1:* 00000000.00000000.0000.00:00:00:00:00:00 0 0 0 0 0 0 0 0 0 0: 0: 0: 5 0:
2 1 1:* 00000000.00000000.0000.00:00:00:00:00:00 0 0 0 0 0 0 0 0 0 0: 0: 0: 1
          17:.../m_dlau/LATEST
3 1 1:* 00000000.00000000.0000.00:00:00:00:00:00 0 0 0 0 0 0 0 1 0 0: 6:m_dlau
0: 1      12:/main/LATEST
4 11 0: 00000000.00000000.0000.00:00:00:00:00:00 0 0 18:/TeamSCM/ClearCase
0:
5 11 0: 00000000.00000000.0000.00:00:00:00:00:00 0 0 14:/TeamSCM/Admin 0:
6 11 0: 00000000.00000000.0000.00:00:00:00:00:00 0 0 22:/ICE/ICEConfig/scripts
0:
```

Reassembling the .compiled_spec file (2 / 5)

- Human readable text file, but more complicated format
- Three main portions
  - Header
  - Element selection rules
  - Load rules

Reassembling the .compiled_spec file (3 / 5)

- Header

  12

  dfe36c75.23684e37.ab97.9f:1f:b4:6d:7c:97

  6

- Note:
  - Retain Line 1 from the parent view
  - Line 2 is the view's uuid – use parent view's uuid
  - Line 3 is the total number of element selection + load rule lines

Reassembling the .compiled_spec file (4 / 5)

- Element selection rules

  1 1 1:* 00000000.00000000.0000.00:00:00:00:00:00 0 0 0 0 0 0 0 0 0 0: 0: 0: 5 0:

  2 1 1:* 00000000.00000000.0000.00:00:00:00:00:00 0 0 0 0 0 0 0 0 0 0: 0: 0: 1
  17:.../m_dlau/LATEST

  3 1 1:* 00000000.00000000.0000.00:00:00:00:00:00 0 0 0 0 0 0 0 1 0 0: 6:m_dlau 0: 1
  12:/main/LATEST

- Corresponding to these lines

  element * CHECKEDOUT

  element * .../m_dlau/LATEST

  element * /main/LATEST -mkbranch m_dlau

- These lines should be the same in every child view

Reassembling the .compiled_spec file (5 / 5)

- Load rules

  4 11 0: 00000000.00000000.0000.00:00:00:00:00:00 0 0 18:/TeamSCM/ClearCase 0:

  5 11 0: 00000000.00000000.0000.00:00:00:00:00:00 0 0 14:/TeamSCM/Admin 0:

  6 11 0: 00000000.00000000.0000.00:00:00:00:00:00 0 0 22:/ICE/ICEConfig/scripts 0:

- Corresponding to these lines

  load /TeamSCM/ClearCase

  load /TeamSCM/Admin

  load /ICE/ICEConfig/scripts

- Notes
  - Copy the load rule line from child to parent
  - First number of each line is the line number
  - Remember to change the line number after copying!

Reassembling the db directory – view database (1 / 3)

- Run “cleartool reformatview –dump” on each view (child and parent)
  - Creates view_db.dump_file in db directory

- Sample view_db.dump_file

```
Ver 5 0 4d911418
FreezeState 1 1
Wso dde39297.34c811d9.8094.00:01:83:08:d6:9e
      ae00b94a.49ee4d03.86e0.ad:9f:54:2f:2f:52
      41:TeamSCM\ClearCase\Tools\crpview02\clients ... <snip>
Wso ...
Wso ...
...<etc>...
Eof
```

Reassembling the db directory – view database (2 / 3)

Wso lines look something like this:

```
Wso dde39297.34c811d9.8094.00:01:83:08:d6:9e
  ae00b94a.49ee4d03.86e0.ad:9f:54:2f:2f:52
  41:TeamSCM\ClearCase\Tools\crpview02\clients
  28257433.276940d7.937f.1f:34:17:8d:e9:44
  f33f38a8.bea44e36.9bfc.a7:d9:4a:89:89:20 8:index.pl
  00000000.00000000.0000.00:00:00:00:00:00
  00000000.00000000.0000.00:00:00:00:00:00
  00000000.00000000.0000.00:00:00:00:00:00 170b 1 4b95a2f5 0 0 3 8f
```

Second last number of each Wso line analogous “Rule: ...” part of “cleartool ls”

element * CHECKEDOUT

element * .../m_dlau/LATEST

element * /main/LATEST -mkbranch m_dlau

Reassembling the db directory – view database (3 / 3)

- Reassemble db file
  - Create the headers (“Ver...” and “FreezeState...” lines) by copying from one of the views
  - Concatenate all “Wso...” lines from child view dump file to parent’s dump file
  - Append “Eof” marker
- Reload database by running “cleartool reformatview –load” on the parent view
- Parent db now contains loaded file info for all files loaded in all child views
- Remove child views when done (cleartool rmview -force)

eBay's Use Case

- Standard developer desktop
  - 2 dual core CPUs (Xeon)
  - 8GB RAM
  - 100Mb ethernet
- View workspace size ~ 5GB
- Across ~100 load rules / ~30 VOBs
- Takes about 45 minutes to create a snapshot view from scratch

Performance Metrics

- Case #1 – default single threaded snapshot view load, using local SATA disk
  - Throughput = 2 MBps , Time = 45 minutes
  - Bottleneck? Sequential single-threaded design
  - Solution? Implement parallel design
- Case #2 – parallel snapshot view load, using local SATA disk
  - Throughput = 5 MBps, Time = 20 minutes
  - Bottleneck? Disk IO
  - Solution? Use SSD drives.
- Case #3 – parallel snapshot view load, using SSD
  - Throughput = 8 MBps, Time = 10 minutes
  - Bottleneck? Possibly network
  - Solution? Work in progress...

Extending the Idea

- Similar technique can be used to replicate entire views, but NOT view privates
- Use Case
  - Lead dev loads his snapshot, does a verification build, rebuild jars etc
  - Everyone else uses a clone of his view as a starting point for development.
- BEFORE (loading view using SSD disk)
  - Time = 10 minutes to load view, 60+ minutes to do initial build
- AFTER (cloning the view)
  - Time = 12 minutes

Considerations and Other Thoughts

- Assumes no symbolic links and non-UCM config spec
- Your Mileage May (WILL!) Vary ☺

www.ibm/software/rational

https://www14.software.ibm.com/iwm/web/cc/earlyprograms/rational.shtml

www.ibm/software/rational

https://www14.software.ibm.com/iwm/web/cc/earlyprograms/rational.shtml